

Python Code

```
from transformers import AutoModelForQuestionAnswering, AutoTokenizer, Trainer, TrainingArguments,
TrainerCallback
from datasets import load_dataset
from pprint import pprint
from tqdm import tqdm
import evaluate
import torch
import logging
import matplotlib.pyplot as plt
import numpy as np

debugging = False

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Device: {device}")

loss_values = [] # Loss function
epoch_progress = []
exact_match_before = 0 # Exact Match score before distillation
exact_match_after = 0 # Exact Match score after distillation
f1_score_before = 0 # F1 score before distillation
f1_score_after = 0 # F1 score after distillation

# Load the SQuAD dataset
squad = load_dataset("squad")

# Load teacher model and tokenizer
#teacher_model_name = "csarron/roberta-base-squad-v1" # First (MVP) teacher model
teacher_model_name = "Dingyun-Huang/roberta-large-squad1" # Best performing teacher model
teacher_model = AutoModelForQuestionAnswering.from_pretrained(teacher_model_name)
teacher_tokenizer = AutoTokenizer.from_pretrained(teacher_model_name)

# Load student model and tokenizer
#student_model_name = "distilroberta-base" # First (MVP) and second model
#student_model_name = "csarron/roberta-base-squad-v1" # Third model
student_model_name = "deepset/roberta-base-squad2" # Fourth and best performing model
student_model = AutoModelForQuestionAnswering.from_pretrained(student_model_name)
student_tokenizer = AutoTokenizer.from_pretrained(student_model_name)

# Define evaluation function
def evaluate_model(model, tokenizer, dataset):
    model.to(device)
    metric = evaluate.load("squad")
    predictions = []
    references = []

    for example in tqdm(dataset, desc="Evaluating"):
        # Tokenize inputs
        inputs = tokenizer(
            example["context"], example["question"], truncation=True,
            padding=True, return_tensors="pt"
        )

        # Move inputs to the same device as the model
        inputs = {key: value.to(device) for key, value in inputs.items()}

        # Get model outputs
        outputs = model(**inputs)
        start_logits, end_logits = outputs.start_logits, outputs.end_logits
        start_idx = torch.argmax(start_logits, dim=-1).item()
        end_idx = torch.argmax(end_logits, dim=-1).item()

        # Decode prediction
        prediction = tokenizer.decode(inputs["input_ids"][0][start_idx:end_idx + 1])

        # Append to predictions
```

```

        predictions.append({
            "id": example["id"],
            "prediction_text": prediction
        })

    # Append to references (ground truth)
    references.append({
        "id": example["id"],
        "answers": example["answers"]
    })

# Compute metrics
result = metric.compute(predictions=predictions, references=references)

return result

# Prepare data for distillation
def preprocess_data(example):
    # Move teacher model to appropriate device
    teacher_model.to(device)

    # Tokenize context and question
    inputs = teacher_tokenizer(
        example["context"],
        example["question"],
        truncation=True,
        padding="max_length",
        max_length=384,
        return_tensors="pt"
    )

    # Move inputs to the same device as the teacher model
    inputs = {key: value.to(device) for key, value in inputs.items()}

    # Get logits from the teacher model
    with torch.no_grad():
        outputs = teacher_model(**inputs)

    # Add teacher logits to the example
    example["input_ids"] = inputs["input_ids"][0].cpu().tolist()
    example["attention_mask"] = inputs["attention_mask"][0].cpu().tolist()
    example["start_logits"] = outputs.start_logits[0].cpu().tolist()
    example["end_logits"] = outputs.end_logits[0].cpu().tolist()

    return example

# Prepare validation data for evaluation
def preprocess_validation_data(example):
    # Tokenize context and question
    inputs = student_tokenizer(
        example["context"],
        example["question"],
        truncation=True,
        padding="max_length",
        max_length=384,
    )

    return inputs

# Evaluate teacher model
print("Teacher Model Evaluation")
result = evaluate_model(teacher_model, teacher_tokenizer, validation_dataset)
print(f"Exact Match: {result['exact_match']:.2f}%")
print(f"F1 Score: {result['f1']:.2f}%\n")

# Evaluate student model on validation set (before distillation)
print("Student Model Evaluation (Before Distillation)")
result = evaluate_model(student_model, student_tokenizer, validation_dataset)
print(f"Exact Match: {result['exact_match']:.2f}%")
print(f"F1 Score: {result['f1']:.2f}%\n")

# Save the results for use in results Graphs

```

```

exact_match_before = result['exact_match']
f1_score_before = result['f1']

# Apply preprocessing to training dataset
if debugging:
    train_dataset = squad["train"].select(range(10000)).map(preprocess_data)
else:
    train_dataset = squad["train"].map(preprocess_data)

# Apply preprocessing to validation dataset
validation_dataset = squad["validation"]
validation_set = validation_dataset.map(preprocess_validation_data, batched=True)

# Define training arguments for student model
training_args = TrainingArguments(
    output_dir="./trained_student",
    logging_steps=10, # Log every 10 steps
    learning_rate=2e-5,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=64,
    num_train_epochs=12,
    weight_decay=0.01,
    save_total_limit=2,
    save_strategy="epoch",
    remove_unused_columns=False,
)

# Custom progress callback used for loss function plot
class TeacherTrainingProgress(TrainerCallback):
    def on_log(self, args, state, control, logs=None, **kwargs):
        if logs and 'loss' in logs:
            loss_values.append(logs['loss'])
            epoch_progress.append(logs['epoch'])

# Define custom loss function for knowledge distillation
class DistillationTrainer(Trainer):
    def compute_loss(self, model, inputs, return_outputs=False, **kwargs):

        # Forward pass
        outputs = model(input_ids=inputs["input_ids"], attention_mask=inputs["attention_mask"])

        # Compute KL Divergence loss between teacher and student logits
        start_loss = torch.nn.functional.kl_div(
            torch.nn.functional.log_softmax(outputs.start_logits, dim=-1),
            torch.nn.functional.softmax(inputs["start_logits"], dim=-1),
            reduction="batchmean"
        )

        end_loss = torch.nn.functional.kl_div(
            torch.nn.functional.log_softmax(outputs.end_logits, dim=-1),
            torch.nn.functional.softmax(inputs["end_logits"], dim=-1),
            reduction="batchmean"
        )

        # Average the start and end losses
        loss = (start_loss + end_loss) / 2

        return (loss, outputs) if return_outputs else loss

class FullyCustomDataCollator:
    def __init__(self, tokenizer):
        self.tokenizer = tokenizer

    def __call__(self, features):
        # Filter out non-tokenized fields before padding
        tokenized_features = [
            {k: v for k, v in f.items() if k in ["input_ids", "attention_mask", "token_type_ids"]}
            for f in features
        ]

        # Dynamically pad input_ids and attention_mask using the tokenizer
        batch = self.tokenizer.pad(

```

```

        tokenized_features,
        padding=True,
        max_length=None,
        return_tensors="pt",
    )

    # Add custom fields (e.g., start_logits and end_logits) to the batch
    if "start_logits" in features[0]:
        batch["start_logits"] = torch.tensor(
            [f["start_logits"] for f in features],
            dtype=torch.float32
        )
    if "end_logits" in features[0]:
        batch["end_logits"] = torch.tensor(
            [f["end_logits"] for f in features],
            dtype=torch.float32
        )

    return batch

```

```

data_collator = FullyCustomDataCollator(tokenizer=student_tokenizer)

```

```

# Train the student model using the knowledge distillation trainer

```

```

trainer = DistillationTrainer(
    model=student_model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=validation_set,
    data_collator=data_collator,
    callbacks=[TeacherTrainingProgress()]
)
train_result = trainer.train()

```

```

# Evaluate student model on validation set after distillation

```

```

print("Student Model Evaluation after Distillation")
result = evaluate_model(student_model, student_tokenizer, validation_dataset)
print(f"Exact Match: {result['exact_match']:.2f}%")
print(f"F1 Score: {result['f1']:.2f}%\n")

```

```

# Save the results for use in Graphs

```

```

exact_match_after = result['exact_match']
f1_score_after = result['f1']

```

```

# Plot loss values per epoch

```

```

plt.figure(figsize=(8, 6))
plt.plot(epoch_progress, loss_values, label="Training Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Loss Per Epoch")
plt.legend()
plt.grid()
plt.show()

```

```

# Graph Exact Match scores before and after distillation

```

```

plt.figure(figsize=(8, 6))
labels = ["Before Distillation", "After Distillation"]
exact_match_scores = [exact_match_before, exact_match_after]
plt.bar(labels, exact_match_scores, color=['blue', 'green'])
plt.ylabel("Exact Match (%)")
plt.title("Exact Match Score Before and After Distillation")
for i, v in enumerate(exact_match_scores):
    plt.text(i, v + 1, f"{v:.1f}%", ha='center', fontsize=12)
plt.show()

```

```

# Graph F1 scores before and after distillation

```

```

plt.figure(figsize=(8, 6))
f1_scores = [f1_score_before, f1_score_after]
plt.bar(labels, f1_scores, color=['orange', 'red'])
plt.ylabel("F1 Score (%)")
plt.title("F1 Score Before and After Distillation")
for i, v in enumerate(f1_scores):
    plt.text(i, v + 1, f"{v:.1f}%", ha='center', fontsize=12)

```

```
plt.show()
```

```
# Save the fine-tuned model and tokenizer to permnament storage
student_model.save_pretrained("./distilled_model")
student_tokenizer.save_pretrained("./distilled_model")
```